


## Neural Networks

161.326

1.1

Stephen Marsland

## McCulloch and Pitts Neurons



- Greatly simplified biological neurons
- Sum the inputs
  - ❖ If total is less than some threshold, neuron fires
  - ❖ Otherwise does not

161.326

1.2

Stephen Marsland

## McCulloch and Pitts Neurons

$$h = \sum_{i=1}^m x_i w_i \quad o = \begin{cases} 1 & h \geq \theta \\ 0 & h < \theta \end{cases} \quad \text{for some threshold } \theta$$

- The weight  $w_j$  can be positive or negative
  - ❖ Inhibitory or excitatory
- Use only a linear sum of inputs
- No refractory period
- Use a simple output instead of a pulse (spike train)

161.326

1.3

Stephen Marsland

## Neural Networks

- Can put lots of McCulloch & Pitts neurons together
- Connect them up in any way we like
- In fact, assemblies of the neurons are capable of *universal computation*
  - ❖ Can perform any computation that a normal computer can
  - ❖ Just have to solve for all the weights  $w_{ij}$

161.326

1.4

Stephen Marsland

## Training Neurons


- Adapting the weights is learning
  - ❖ How does the network know it is right?
  - ❖ How do we adapt the weights to make the network right more often?
- Training set with target outputs
- Learning rule

161.326

1.5

Stephen Marsland

## The Perceptron Network



161.326

1.6

Stephen Marsland

## Updating the Weights

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$

- We want to change the values of the weights
- Aim: minimise the *error* at the output
- If  $E = t - y$ , want  $E$  to be 0
- Use:

$$\Delta w_{ij} = \eta \cdot (t_j - y_j) \cdot x_i$$


Learning rate  $\eta$       Input  $x_i$       Error  $(t_j - y_j)$

161.326

1.7

Stephen Marsland

## Biases Replace Thresholds

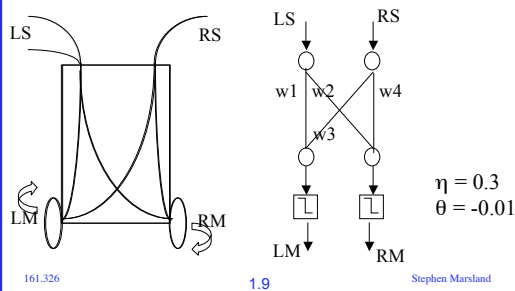


161.326

1.8

Stephen Marsland

## Obstacle Avoidance with the Perceptron



161.326

1.9

Stephen Marsland

## Obstacle Avoidance with the Perceptron


LS	RS	LM	RM
0	0	1	1
0	1	-1	1
1	0	1	-1
1	1	X	X

161.326

1.10

Stephen Marsland

## Obstacle Avoidance with the Perceptron



$$\Delta w_{ij} = \eta \cdot (t_j - y_j) \cdot x_i$$


$$w1 = 0 + 0.3 \cdot (1 - 1) \cdot 0 = 0$$

161.326

1.11

Stephen Marsland

## Obstacle Avoidance with the Perceptron



$$w2 = 0 + 0.3 \cdot (1 - 1) \cdot 0 = 0$$

And the same for w3, w4

161.326

1.12

Stephen Marsland

## Obstacle Avoidance with the Perceptron


LS	RS	LM	RM
0	0	1	1
0	1	-1	1
1	0	1	-1
1	1	X	X

161.326

1.13

Stephen Marsland

## Obstacle Avoidance with the Perceptron




$$w1 = 0 + 0.3 * (-1 - 1) * 0 = 0$$

161.326

1.14

Stephen Marsland

## Obstacle Avoidance with the Perceptron



$$w1 = 0 + 0.3 * (-1 - 1) * 0 = 0$$


$$w2 = 0 + 0.3 * (1 - 1) * 0 = 0$$

161.326

1.15

Stephen Marsland

## Obstacle Avoidance with the Perceptron



$$w1 = 0 + 0.3 * (-1 - 1) * 0 = 0$$

$$w2 = 0 + 0.3 * (1 - 1) * 0 = 0$$


$$w3 = 0 + 0.3 * (-1 - 1) * 1 = -0.6$$

161.326

1.16

Stephen Marsland

## Obstacle Avoidance with the Perceptron



$$w1 = 0 + 0.3 * (-1 - 1) * 0 = 0$$

$$w2 = 0 + 0.3 * (1 - 1) * 0 = 0$$

$$w3 = 0 + 0.3 * (-1 - 1) * 1 = -0.6$$

$$w4 = 0 + 0.3 * (1 - 1) * 1 = 0$$

161.326

1.17

Stephen Marsland

## Obstacle Avoidance with the Perceptron


LS	RS	LM	RM
0	0	1	1
0	1	-1	1
1	0	1	-1
1	1	X	X

161.326

1.18

Stephen Marsland

## Obstacle Avoidance with the Perceptron




$$\begin{aligned} w1 &= 0 + 0.3 * (1-1) * 1 = 0 \\ w2 &= 0 + 0.3 * (-1-1) * 1 = -0.6 \\ w3 &= -0.6 + 0.3 * (1-1) * 0 = -0.6 \\ w4 &= 0 + 0.3 * (-1-1) * 0 = 0 \end{aligned}$$

161.326

1.19

Stephen Marsland

## Obstacle Avoidance with the Perceptron



161.326

1.20

Stephen Marsland

## Implementation

- The forward stage of the network is easy:

```
activations = dot(inputs,weights)
activations = where(activations>0,1,0)
```

161.326

1.21

Stephen Marsland

## Implementation

- Need to include the bias node bit:

```
inputs = concatenate((ones
((self.nData,1)),inputs),axis=1)
```

161.326

1.22

Stephen Marsland

## Implementation

- And then it is just the weight update:

```
weights += eta*dot(transpose(inputs),
targets-activations)
```

That's pretty much all there is to it

161.326

1.23

Stephen Marsland

## Linear Separability

- Outputs are:

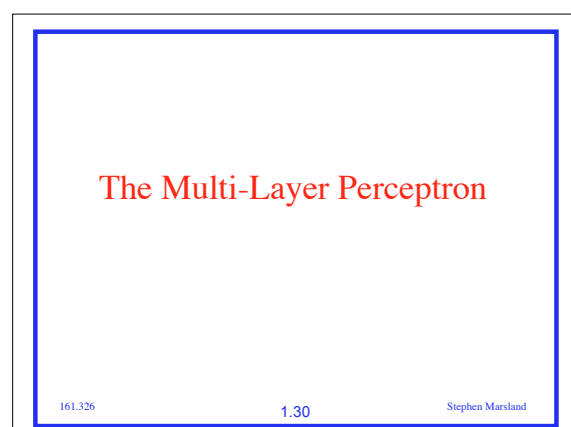





$$y_j = \text{sign} \left( \sum_{i=1}^n w_{ij} x_i \right)$$


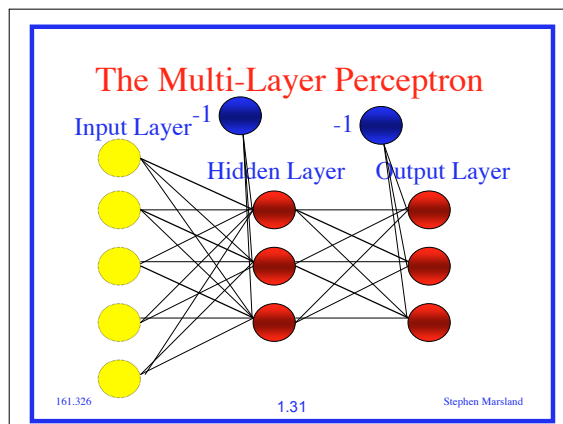
$$\Rightarrow \mathbf{w} \cdot \mathbf{x} > 0$$

161.326

1.24

Stephen Marsland




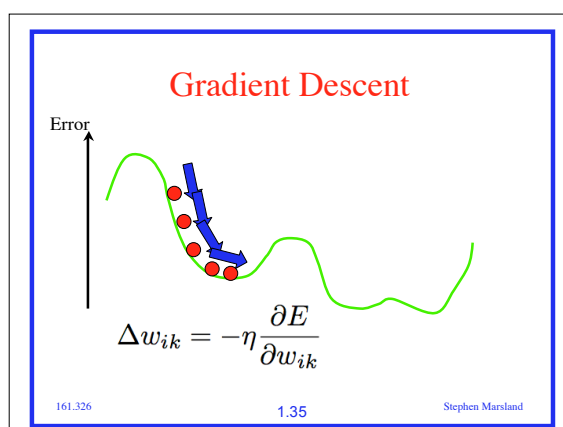


### XOR Again

A	B	C <sub>in</sub>	C <sub>out</sub>	D <sub>in</sub>	D <sub>out</sub>	E
0	0	-0.5	0	-1	0	-0.5
0	1	0.5	1	0	0	0.5
1	0	0.5	1	0	0	0.5
1	1	1.5	1	1	1	-0.5

161.326    1.33    Stephen Marsland

- ### Gradient Descent
- The MLP can solve XOR
  - How do we choose the weights?
  - Harder than Perceptron
    - ❖ More weights
    - ❖ Which weights are wrong? Input-hidden or hidden-output?
  - Use gradient descent learning
  - Compute gradient -> differentiation
- 161.326    1.34    Stephen Marsland



## An Error Function

- For Perceptron, looked at  $(t-y)$
- Better: sum-of-squares error

$$E(\mathbf{w}) = \frac{1}{2} \sum_k (t_k - y_k)^2 = \frac{1}{2} \sum_k \left( t_k - \sum_i w_{ik} x_i \right)^2$$

- One more thing - we will ignore the threshold function in the neurons

$$\Rightarrow \frac{\partial E}{\partial w_{ik}} = \sum_k (t_k - y_k)(-x_i)$$


161.33

Stephen Marsland

## Training MLPs

### (1) Forward Pass

- ❖ Put the input values in the input layer
- ❖ Calculate the activations of the hidden nodes
- ❖ Calculate the activations of the output nodes
- ❖ Calculate the errors using the targets



161.326

1.38

Stephen Marsland

## Training MLPs

- For output nodes
  - ❖ Don't know input
- For hidden nodes
  - ❖ Don't know targets
- For extra hidden layers
  - ❖ Don't know either
- Therefore, hard to use gradient descent

161.326

1.39

Stephen Marsland

## Backpropagation of Error

161.326


1.40

Stephen Marsland

## Training MLPs

### (2) Backward Pass

- ❖ From output errors, update last layer of weights
- ❖ From these errors, update next layer
- ❖ Work backwards through the network
- ❖ Error is backpropagated through the network



161.326

1.41

Stephen Marsland

## Activation Function

- In the analysis we've ignored the activation function
  - ❖ The thresholder is not differentiable
- What do we want in an activation function?
  - ❖ Differentiable
  - ❖ Should saturate (become constant at ends)
  - ❖ Change between saturation values quickly


161.326

1.42

Stephen Marsland

## Sigmoid Functions

$$g(a) = \frac{1}{1 + \exp(-\beta a)}$$



161.326

Stephen Marsland

## Error Terms

- Need to differentiate the sigmoid function
- Gives us the following *error terms* (deltas)

❖ For the outputs

$$\delta_k = (y_k - t_k) y_k (1 - y_k)$$

❖ For the hidden nodes

$$\delta_j = a_j (1 - a_j) \sum_k w_{jk} \delta_k$$

161.326

1.44

Stephen Marsland

## Update Rules

- This gives us the necessary update rules

❖ For the weights connected to the outputs:

$$w_{jk} \leftarrow w_{jk} - \eta \delta_k a_j^{\text{hidden}}$$

❖ For the weights connect to the hidden nodes:

$$v_{ij} \leftarrow v_{ij} - \eta \delta_j x_i$$

161.326

1.45

Stephen Marsland

## Summary of Backpropagation

- Introduce inputs
- Feed values forward through network
- Compute sum-of-squares error at outputs
- Compute the delta terms at the output by differentiation
- Use this to update the weights from the outputs to the last hidden layer

161.326

1.46

Stephen Marsland

## Summary of Backpropagation

- Once these are correct, propagate errors back to the neurons of the hidden layers
- Compute the delta terms for these neurons
- Use them to update the next set of weights
- Repeat until reach the inputs

161.326

1.47

Stephen Marsland

## Implementation

- Forwards isn't much different to the Perceptron (except do it twice):

```
inputs = concatenate((inputs, -ones
((self.ndata, 1))), axis=1)
hidden = dot(inputs, weights1);
hidden = 1.0 / (1.0 + exp(-beta * hidden))
hidden = concatenate((hidden, -ones
((ndata, 1))), axis=1)
outputs = dot(hidden, weights2);
return 1.0 / (1.0 + exp(-beta * outputs))
```

161.326

1.48

Stephen Marsland



## Implementation

- The updates are more involved - here's the one for the output weights

```
deltah = zeros(nhidden+1)
for j in range(nhidden+1):
    sumk = sum(weights2[j,:]*deltao[d,:])
    deltah[j] = hidden[d,j]*
        (1.0-hidden[d,j])*sumk
```

161.326

1.49

Stephen Marsland

## Implementation

- One new function

```
random.shuffle(change)
inputs = inputs[change,: ]
targets = targets[change,: ]
```

161.326

1.50

Stephen Marsland

## Network Topology

- How many layers?
- How many neurons per layer?
- No good answers
  - ❖ At most 3 layers, usually 2
  - ❖ Guess size of layers (usually get smaller)
  - ❖ Test several different networks

161.326

1.51

Stephen Marsland

## Batch Learning

- When should the weights be updated?
  - ❖ After all inputs seen (batch)
    - ✓ More accurate estimate of gradient
    - ✓ Converges to local minimum faster
  - ❖ After each input is seen (sequential)
    - ✓ Simpler to program
    - ✓ May escape from local minima (change order or presentation)
- Both ways, need many epochs - passes through the whole dataset

161.326

1.52

Stephen Marsland

## Momentum



$$w_{ij}^{\tau} \leftarrow w_{ij}^{\tau-1} \eta \delta_j^{\text{hidden}} + \alpha \Delta w_{ij}^{\tau-1},$$


- Give more weight to the ball
- Can use smaller learning rate (more stable)
- May overcome local minima

161.326

1.53

Stephen Marsland

## Learning Capacity



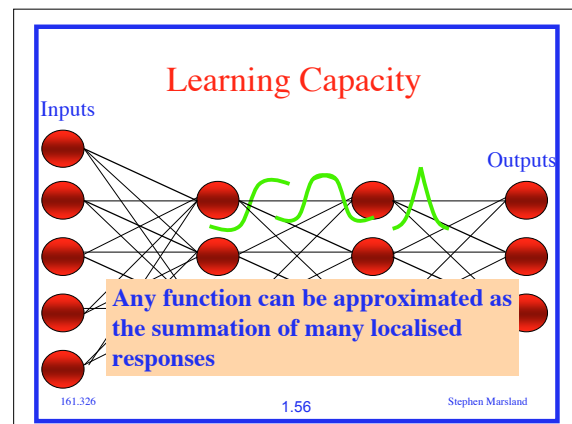
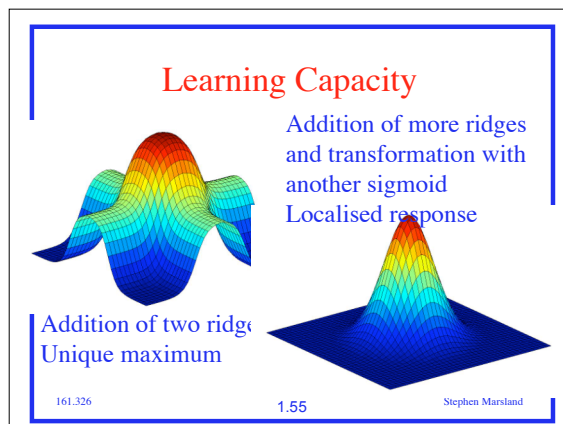
Output of one sigmoid

Addition of two sigmoids

161.326

1.54

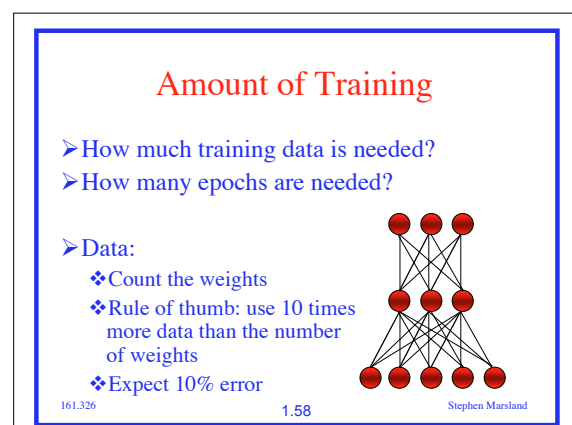
Stephen Marsland



## Decision Boundaries

Structure	Types of Decision Regions	Exclusive OR Problem	Classes with Meshed Regions	Most General Region Shapes
Single-Layer	Half Plane Bounded by Hyperplane			
Two-Layer	Convex Open or Closed Regions			
Three-Layer	Arbitrary (Complexity Limited by Number of Nodes)			

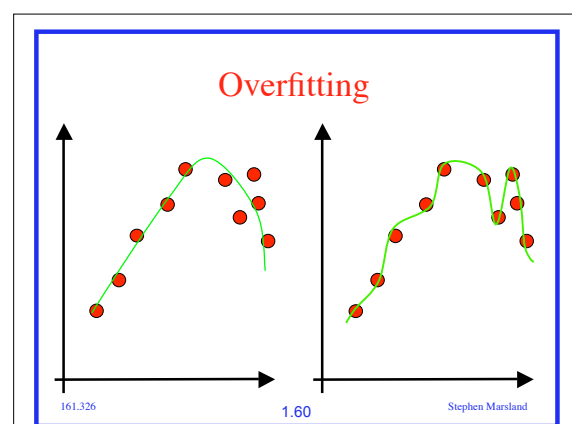
161.326 1.57 Stephen Marsland



## Generalisation

- Aim of neural network learning:
- Generalise from training examples to all possible inputs
- Undertraining is bad
- Overtraining is worse

161.326 1.59 Stephen Marsland



## Overfitting

- MLP has easily enough variation to fit any surface
- We want to learn the data without the noise
- Overtraining lets the network overfit
  - ❖ Then does not generalise
  - ❖ Function is too complicated

161.326

1.61

Stephen Marsland

## Testing

- How do we evaluate our trained network?
- Can't just compute the error on the training data - unfair, can't see overfitting
- Keep a separate testing set
- After training, evaluate on this test set
- How do we check for overfitting?
- Can't use training or testing sets

161.326

1.62

Stephen Marsland

## Validation


- Keep a third set of data for this
- Train the network on training data
- Periodically, stop and evaluate on validation set
- After training has finished, test on test set
- This is coming expensive on data!

161.326

1.63

Stephen Marsland

## Hold Out Cross Validation



161.326

1.64

Stephen Marsland

## Hold Out Cross Validation

- Partition training data into K subsets
- Train on K-1 of subsets, validate on Kth
- Repeat for new network, leaving out a different subset
- Choose network that has best validation error
- Traded off data for computation
- Extreme version: leave-one-out

161.326

1.65

Stephen Marsland



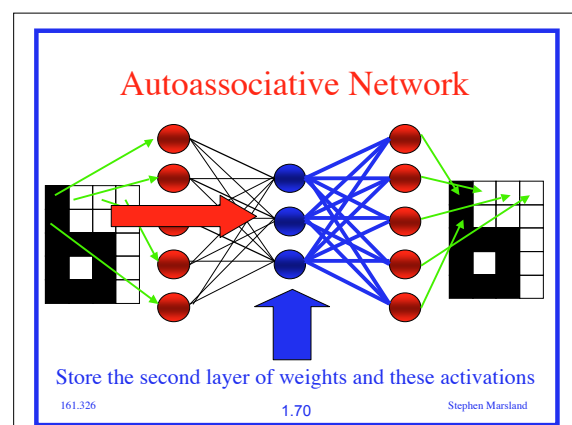
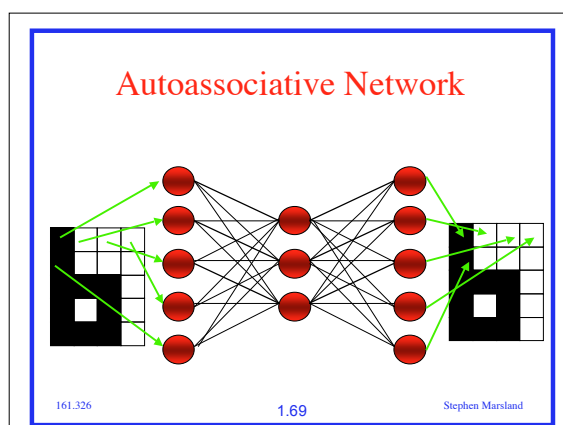
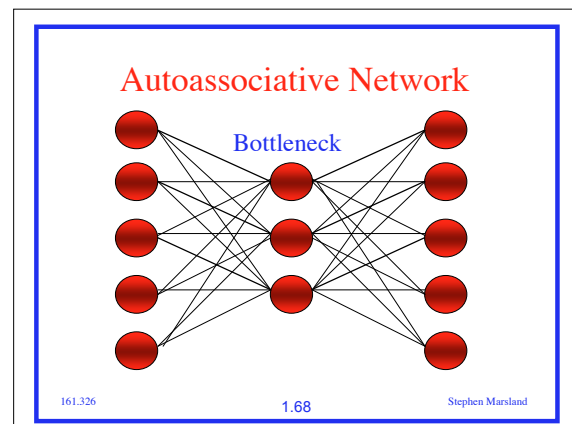
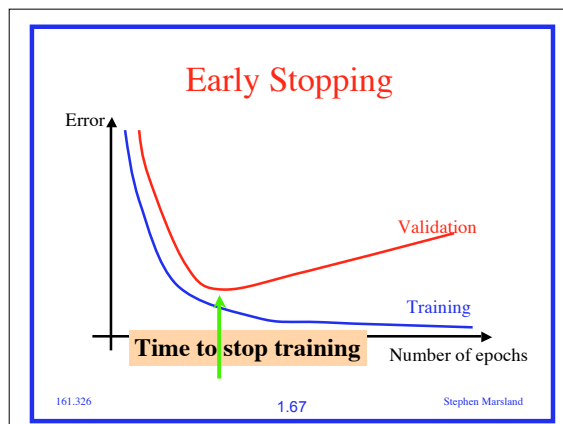
## Early Stopping

- When should we stop training?
  - ❖ Could set a minimum training error
    - ✓ Danger of overfitting
  - ❖ Could set a number of epochs
    - ✓ Danger of underfitting or overfitting
  - ❖ Can use the validation set
    - ✓ Measure the error on the validation set during training

161.326

1.66

Stephen Marsland



## Classification

161.326

1.73

Stephen Marsland

## What is Classification?

- Based on exemplars of each class, choose which of  $N$  classes an input belongs to
- Split the input set into the  $N$  classes
- What about inputs that aren't in any of the classes?
  - ❖ Novelty detection

161.326

1.74

Stephen Marsland

## Example: Vending Machine

- Recognise coins fed into machine
- Choose features
  - ❖ Weight
  - ❖ Diameter
  - ❖ Shape
  - ❖ Colour



161.326

1.75

Stephen Marsland

## Example: Vending Machine

- Based on exemplars of each class, calculate feature values
- Use these to recognise coins
- What about foreign coins?
  - ❖ Novelty detection




161.326

1.76

Stephen Marsland

## Decision Boundaries




161.326

1.77

Stephen Marsland

## Decision Boundaries



161.326

1.78

Stephen Marsland

## Choosing Features

- Art or science?
  - ❖ Not too many features - curse of dimensionality
  - ❖ Enough, so can separate classes
  - ❖ Not all features are equally useful

161.326

1.79

Stephen Marsland

## Classification with the Multi-Layer Perceptron

161.326

1.80

Stephen Marsland

## Classification & the MLP

- Inputs to MLP are feature values
- What about outputs?
- Use one neuron with:

$$\left\{ \begin{array}{ll} C_1 & \text{if } y \leq -0.5 \\ C_2 & \text{if } -0.5 < y \leq 0 \\ C_3 & \text{if } 0 < y \leq 0.5 \\ C_4 & \text{if } y > 0.5 \end{array} \right\}$$

161.326

1.81

Stephen Marsland

## 1 of N Encoding

- 1 output neuron for each class
- Target data has 1 output true, rest false
- In use, pick class as the neuron with the highest activation
- Often use softmax activation for output nodes

$$\frac{\exp(z_i)}{\sum_{j=1}^N \exp(z_j)}$$

161.326

1.82

Stephen Marsland