

Optimisation

- Most of this course has been about optimisation
 - ❖ Gradient descent
- What if there is no gradient?
 - ❖ Discrete problems
 - ❖ Check all cases?

1.59.302

5.2

Stephen Marsland

No Free Lunch

- There is no perfect search algorithm
- You always need to think about the problem
- Might have to design the encoding of the problem carefully
- In practice, gradient descent is pretty good for continuous problems

1.59.302

5.4

Stephen Marsland

Search and Genetic Algorithms

Stephen Marsland


1.59.302

5.1

Stephen Marsland

Discrete Problems

- Chip Design
 - ❖ Position circuits on a chip so that none of the lines cross
- Timetabling
 - ❖ Given courses and students, find a timetable with the minimum number of clashes



1.59.302

5.3

Stephen Marsland

The Travelling Salesman

- An NP-complete problem
- There are $N!$ different possible solutions (where N is the number of cities)
 - ❖ Virtually impossible to solve for $N > 10$
- Actually useful - logistics, chip design

1.59.302

5.6

Stephen Marsland

The Travelling Salesman

Find a tour that starts and ends at the same city, visits every city precisely once, and has the minimum possible distance



1.59.302

5.5

Stephen Marsland

Greedy Search

- Choose a start city
- Repeat
 - ❖ Pick the closest remaining city
- Until get back to start city
- No backtracking

1.59.302

5.8

Stephen Marsland

Exhaustive Search

- Try out every solution
- Trivial to implement
- $O(N!)$ worse than $O(\exp(N))$

1.59.302

5.7

Stephen Marsland

Hill Climbing

- Pick a starting cycle
- Repeat
 - ❖ Swap a random pair of cities
 - ❖ If tour length decreases keep the new order
- Until you get bored or stops changing

1.59.302

5.10

Stephen Marsland

Greedy Search

- Computationally cheap: $O(N)$
- Will find a solution
- No guarantee of any kind of optimality
- Cannot predict how good solution is

1.59.302

5.9

Stephen Marsland

Hill Climbing


- Foothills
- ❖ Local maxima
- Plateaus
- ❖ Algorithm fails to get anywhere at all
- Ridges
- ❖ Look like peaks - most directions go down

1.59.302

5.12

Stephen Marsland

Hill Climbing




1.59.302

5.11

Stephen Marsland

Exploitation and Exploration

- Room of 1-armed bandits
- Maximise your payout
- Do you
 - ❖ Use the best machine you've found so far
 - ❖ Try out some more hoping to find a better one



1.59.302

5.14

Stephen Marsland

Exploitation and Exploration

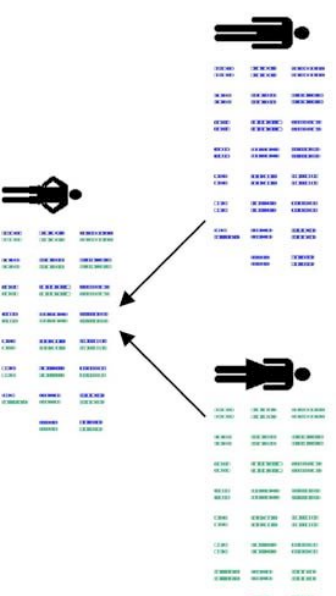
- Try out new solutions
 - ❖ Exploration
 - ❖ Exhaustive search
- Try to improve your current best solution
 - ❖ Exploitation
 - ❖ Hill climbing
- Ideally, some of both

1.59.302

5.13

Stephen Marsland

Evolution



1.59.302

5.16

Stephen Marsland

Evolution

- Survival of the fittest
 - ❖ Live longer (more chance to reproduce)
 - ❖ More attractive (more chance to reproduce)
 - ❖ Higher number of offspring (more survive)
- 50% chance of any gene being in offspring

1.59.302

5.15

Stephen Marsland

Fitness Landscapes


- Useful imagination aid
- Animals fitness is partly due to competition with other animals and environment
- Changes over time
- Evolution favours animals that evolve to peaks of the fitness landscape

1.59.302

5.18

Stephen Marsland

Punctuated Equilibrium



1.59.302

5.17

Stephen Marsland


The Genetic Algorithm

1.59.302

5.20

Stephen Marsland

Fitness Landscapes



1.59.302

5.19

Stephen Marsland

Representation: Strings and Things

Chromosome	String
Gene	Elements
Allele	Alphabet

1.59.302

5.22

Stephen Marsland

Representation: Strings and Things

- Example: bill paying
- List of 100 bills to pay
- Use string of 100 elements
- Each element is whether to pay one bill
- 10110 means pay bills 1, 3, and 4

1.59.302

5.24

Stephen Marsland

The Genetic Algorithm

- How can we abstract the useful bits of evolution?
- How can we fit them into an algorithm?
- Does it work?

1.59.302

5.21

Stephen Marsland

Representation: Strings and Things

- Work out a way of encoding problem
- Choose an alphabet
 - ❖ Possible values of each element of string
 - ❖ Often binary
- Not always easy
- Split up the problem into discrete parts

1.59.302

5.23

Stephen Marsland

Selection

- Choosing parents is crucial
- Want the best (fittest) strings to reproduce
 - ❖ Exploitation
- What about non-fit strings?
 - ❖ Exploration
- Generate a 'mating pool'

1.59.302

5.26

Stephen Marsland

Fitness Functions

- Decide how good the string is
- You pass in a string and get back a number
- The higher the number, the better the solution
- Generally a positive number
- Problem-specific part


1.59.302

5.25

Stephen Marsland

Selection Methods

- Like a roulette wheel
- Probability of picking a string is proportional to its area on wheel
- Fitter strings have a larger number of copies - larger area



1.59.302

5.28

Stephen Marsland

Selection Methods

- If use uniform distribution, all strings equally likely
- Probably not good
- Fitness proportional selection
 - ❖ Pick proportional to fitness

$$p^{\alpha} = \frac{F^{\alpha}}{\sum_{\alpha'} F^{\alpha'}}$$

1.59.302

5.27

Stephen Marsland

Selection Methods

- Niching
 - ❖ Sometime exploration stops - premature convergence
 - ❖ Evolve several subpopulations and occasionally swap a few elements
 - ❖ Also called island populations

1.59.302

5.30

Stephen Marsland

Selection Methods

- Truncation selection
 - ❖ Pick the top 50% of strings
 - ❖ Choose from them at random
- Elitism
 - ❖ Keep a copy of the best strings all the time
- Tournaments
 - ❖ Put the fittest 2 out of the parents and offspring into the new population

1.59.302

5.29

Stephen Marsland

Genetic Operators

➤ Single Point Crossover

```
10011000101
01111010110
-----
10011010110
01111000101
```

1.59.302

5.32

Stephen Marsland

Genetic Operators

➤ How do we combine the two parents?

```
10011000101
01111010110
```

1.59.302

5.31

Stephen Marsland

Genetic Operators

➤ Uniform Crossover

Random Samples

0 0 1 1 0 1 1 0 1 1 0

String 0 1 0 0 1 1 0 0 0 1 0 1

String 1 0 1 1 1 1 0 1 0 1 1 0

1 0 1 1 1 0 1 0 1 1 1

159,302

5.34

Stephen Marsland

Putting It All Together

- Generate a random population of strings
- Repeat
 - ❖ Compute their fitnesses
 - ❖ Select the mating pool
 - ❖ Crossover parents to produce offspring
 - ❖ Mutate the results
 - ❖ Put them in the new population
- Until stopping criteria met

159,302

5.36

Stephen Marsland

Genetic Operators

➤ Multi-Point Crossover

1 0 0 1 1 0 0 0 1 0 1
0 1 1 1 1 0 1 0 1 1 0

1 0 0 1 1 0 1 0 1 0 1

159,302

5.33

Stephen Marsland

Genetic Operators

➤ Mutation

1 0 1 1 1 0 1 0 1 1 1
1 0 1 1 0 0 1 0 1 1 1




159,302

5.35

Stephen Marsland

The Problem



- Colour the map with any of three colours
- No two bordering countries can have the same colour


1.59.302

5.38

Stephen Marsland

Making a GA for 3-Colouring

- Devise a fitness function
 - ❖ Fitness increases for better solutions
 - ❖ F = number of boundaries that are OK



16 out of 26 in this example

1.59.302

5.40

Stephen Marsland

Genetic Algorithms in Practise: Colouring A Map

1.59.302


5.37

Stephen Marsland

Making a GA for 3-Colouring

- Represent possible solutions as strings

RGGBRRBGGGBRRB




1.59.302

5.39

Stephen Marsland

Making a GA for 3-Colouring



- Generate a new population from the old one
- Bias the selection to pick the fitter strings more often

1.59.302

5.42

Stephen Marsland

Making a GA for 3-Colouring


- Make a population of solutions
- [R G B B ...] [G G G R ...] [R G R B ...]
 [G B R R ...] [B G R B ...] [G B R G ...]
- Apply genetic operators to them to make a new population
 - Iterate

1.59.302

5.41

Stephen Marsland

Making a GA for 3-Colouring



Crossover


- Combine pairs of strings to form new strings

1.59.302

5.44

Stephen Marsland

Making a GA for 3-Colouring



Mutation


- Change randomly selected bit to another colour with small probability

1.59.302

5.43

Stephen Marsland


Making a GA for 3-Colouring



1.59.302

Stephen Marsland


Making a GA for 3-Colouring



1.59.302

Stephen Marsland

Making a GA for 3-Colouring



1.59.302

Stephen Marsland



How Important Is Crossover?

1.59.302

5.47

Stephen Marsland

Making a GA for 3-Colouring




159,302

5.50

Stephen Marsland


Making a GA for 3-Colouring



159,302

Stephen Marsland


Making a GA for 3-Colouring



159,302

Stephen Marsland

Making a GA for 3-Colouring



159,302

5.51

Stephen Marsland

Premature Convergence

- Fitter members of population are favoured
- Solutions at local maxima are favoured
 - ❖ Exploitation, not exploration
 - ❖ Selection for the local maximum
 - ❖ Diversity in population reduces
- Already seen niching
- Can also average fitness across number of identical strings (fitness sharing)

1.59.302

5.54

Stephen Marsland

Neural Networks and Genetic Algorithms


- Multi-Layer Perceptron
 - ❖ Computed error at each neuron
 - ❖ Computed gradient
- Could use GA instead of gradient descent
 - ❖ Throw away gradient information
 - ❖ Replace all errors by 1 fitness

1.59.302

5.56

Stephen Marsland

Making a GA for 3-Colouring



1.59.302

Stephen Marsland

Guaranteed Convergence?

- Not much successful analysis
 - ❖ Not guaranteed to converge at all
 - ❖ Not guaranteed to reach global maximum
 - ❖ Can be very slow
- Trade-off between the genetic operators

1.59.302

5.55

Stephen Marsland

Genetic Programming

- Represent computer programs as trees
- Evolve the trees
- Fitness is how well/quickly the program works
- Used for many tasks
 - ❖ Skin melanoma detection
 - ❖ Chip design

1.59.302

5.58

Stephen Marsland

Neural Networks and Genetic Algorithms


- More sensible: use GA to select the network structure
- ❖ Mutation only - crossover not sensible
- ❖ Several mutations:
 - ✓ Add a node
 - ✓ Delete a node
 - ✓ Add a weight connection
 - ✓ Delete a weight connection
- ❖ Use normal training

1.59.302

5.57

Stephen Marsland

Mutation



Choose a subtree to replace


1.59.302

5.60

Stephen Marsland

Generate a

random subtree




(+ 1 2 (IF (> TIME 10) 3 4))

1.59.302

5.59

Stephen Marsland

Crossover



Choose a pair of subtrees to swap


1.59.302

5.62

Stephen Marsland

Genetic Programming

4 randomly chosen initial trees




1.59.302

5.64

Stephen Marsland

Mutation




1.59.302

5.61

Stephen Marsland

Crossover



1.59.302

5.63

Stephen Marsland

Genetic Programming


- Search space unbelievably large
- Depends strongly on initial population
- Programmer develops a set of possibly useful subtrees

1.59.302

5.66

Stephen Marsland

Genetic Programming



- Copy of (a)
- Mutation of (c) from the (2)
- Crossover of (a) and (b)
- Second crossover of (a) and (b)

1.59.302

5.65

Stephen Marsland